

AlphaTensor


7 Nov 2022

John Tan Chong Min

Overview

- Neural Networks typically are implemented as Matrix (Tensor) multiplications
- A speedup in matrix multiplication (costly operation is multiplication step) will help to speed up Deep Learning
- AlphaTensor uses Deep Reinforcement Learning (DRL) to find an improvement to reduce multiplication steps in matrix multiplication
 - 47 for 4x4 by 4x4 (previously 49)
 - 96 for 5x5 by 5x5 (previously 98)

How AlphaGo Zero Works






- Train so that policy network mimics the improved policy distribution π generated by MCTS
- Train so that value network output matches the game outcome for all time steps
- Joint Loss function:

$$(\mathbf{p}, v) = f_{\theta}(s) \text{ and } l = (z - v)^2 - \pi^T \log \mathbf{p} + c \|\theta\|^2$$
- z : Actual game outcome, v : Predicted game outcome
- π : MCTS probability of next move, p : Predicted probability of next move
- c : L2 regularization constant (set at 10^{-4})



Policy Network (Breadth)

- Focus on moves with higher win rate!



Value Network (Depth)

- Look ahead without simulation!
- Even more accurate than Monte Carlo rollouts



Neural Network:
-1

AlphaTensor


Single-player AlphaZero applied to Matrix Multiplication

Overview

- Formulate matrix multiplication as a single-player game, TensorGame
- At each step, player selects how to combine different entries of matrices to multiply
- Score assigned based on number of multiplications needed, and complexity of the resultant matrix

The Simplified Game

- You need to turn off all the lights of this row:



- Each turn, you can only turn on/off all lights at the same time of a few sets of bulbs
 - {1, 2}
 - {2, 4}
 - {3, 5}
 - {4, 5}
- How do you turn off all the lights with the least number of moves?

Normal approach to multiplication

- Need 8 product steps

- $c_1 = a_1b_1 + a_2b_3$

- $c_2 = a_1b_2 + a_2b_4$


- $c_3 = a_3b_1 + a_4b_3$

- $c_4 = a_3b_2 + a_4b_4$

Can be viewed as specifying a 3D 4x4x4 tensor
(with entries being 0 and 1) to specify which are the
product terms for resultant matrix C

a

$$\begin{pmatrix} c_1 & c_2 \\ c_3 & c_4 \end{pmatrix} = \begin{pmatrix} a_1 & a_2 \\ a_3 & a_4 \end{pmatrix} \cdot \begin{pmatrix} b_1 & b_2 \\ b_3 & b_4 \end{pmatrix}$$



Strassen's Algorithm

Assuming summation terms have 0 cost, we seek to reduce number of multiplication terms

- Form intermediate terms m based on (sum of selected a factors)(sum of selected b factors)
- Sum up selected m factors to form c
- We can achieve it in 7 multiplication steps!

b

$$m_1 = (a_1 + a_4)(b_1 + b_4)$$

$$m_2 = (a_3 + a_4)b_1$$

$$m_3 = a_1(b_2 - b_4)$$

$$m_4 = a_4(b_3 - b_1)$$

$$m_5 = (a_1 + a_2)b_4$$

$$m_6 = (a_3 - a_1)(b_1 + b_2)$$

$$m_7 = (a_2 - a_4)(b_3 + b_4)$$

$$c_1 = m_1 + m_4 - m_5 + m_7$$

$$c_2 = m_3 + m_5$$

$$c_3 = m_2 + m_4$$

$$c_4 = m_1 - m_2 + m_3 + m_6$$

c

$$\mathbf{U} = \begin{pmatrix} 1 & 0 & 1 & 0 & 1 & -1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & -1 \end{pmatrix}$$

$$\mathbf{V} = \begin{pmatrix} 1 & 1 & 0 & -1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & -1 & 0 & 1 & 0 & 1 \end{pmatrix}$$

$$\mathbf{W} = \begin{pmatrix} 1 & 0 & 0 & 1 & -1 & 0 & 1 \\ 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & -1 & 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$

Main algorithm

Algorithm 1

A meta-algorithm parameterized by $\{\mathbf{u}^{(r)}, \mathbf{v}^{(r)}, \mathbf{w}^{(r)}\}_{r=1}^R$ for computing the matrix product $\mathbf{C} = \mathbf{AB}$. It is noted that R controls the number of multiplications between input matrix entries.

Parameters: $\{\mathbf{u}^{(r)}, \mathbf{v}^{(r)}, \mathbf{w}^{(r)}\}_{r=1}^R$: length- n^2 vectors such that $\mathcal{I}_n = \sum_{r=1}^R \mathbf{u}^{(r)} \otimes \mathbf{v}^{(r)} \otimes \mathbf{w}^{(r)}$

Input: \mathbf{A}, \mathbf{B} : matrices of size $n \times n$

Output: $\mathbf{C} = \mathbf{AB}$

(1) **for** $r=1, \dots, R$ **do**

(2) $m_r \leftarrow (u_1^{(r)}a_1 + \dots + u_{n^2}^{(r)}a_{n^2})(v_1^{(r)}b_1 + \dots + v_{n^2}^{(r)}b_{n^2})$

(3) **for** $i=1, \dots, n^2$ **do**

(4) $c_i \leftarrow w_i^{(1)}m_1 + \dots + w_i^{(R)}m_R$

return \mathbf{C}

$$m_1 = (a_1 + a_4)(b_1 + b_4)$$

$$m_2 = (a_3 + a_4)b_1$$

$$m_3 = a_1(b_2 - b_4)$$

$$m_4 = a_4(b_3 - b_1)$$

$$m_5 = (a_1 + a_2)b_4$$

$$m_6 = (a_3 - a_1)(b_1 + b_2)$$

$$m_7 = (a_2 - a_4)(b_3 + b_4)$$

$$c_1 = m_1 + m_4 - m_5 + m_7$$

$$c_2 = m_3 + m_5$$

$$c_3 = m_2 + m_4$$

$$c_4 = m_1 - m_2 + m_3 + m_6$$

$$\mathbf{u} = \begin{pmatrix} 1 & 0 & 1 & 0 & 1 & -1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & -1 \end{pmatrix}$$

$$\mathbf{v} = \begin{pmatrix} 1 & 1 & 0 & -1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & -1 & 0 & 1 & 0 & 1 \end{pmatrix}$$

$$\mathbf{w} = \begin{pmatrix} 1 & 0 & 0 & 1 & -1 & 0 & 1 \\ 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & -1 & 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$

How to play the game

- Start with the desired 3D tensor for the matrix multiplication

- Each step r , we then subtract the tensor

$$\mathbf{u}^{(r)} \otimes \mathbf{v}^{(r)} \otimes \mathbf{w}^{(r)}$$

- Eventually after R steps, we will have subtracted $\mathcal{T}_n = \sum_{r=1}^R \mathbf{u}^{(r)} \otimes \mathbf{v}^{(r)} \otimes \mathbf{w}^{(r)}$

- If the list of steps are correct, we will get the $\mathbf{0}$ tensor

- Definition of outer product (u is vector of length m , v vector of length n):

$$\mathbf{u} \otimes \mathbf{v} = \mathbf{A} = \begin{bmatrix} u_1 v_1 & u_1 v_2 & \dots & u_1 v_n \\ u_2 v_1 & u_2 v_2 & \dots & u_2 v_n \\ \vdots & \vdots & \ddots & \vdots \\ u_m v_1 & u_m v_2 & \dots & u_m v_n \end{bmatrix}$$

$r = 1$

$$\mathbf{u} = \begin{pmatrix} 1 & 0 & 1 & 0 & 1 & -1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & -1 \end{pmatrix}$$

$$\mathbf{v} = \begin{pmatrix} 1 & 1 & 0 & -1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & -1 & 0 & 1 & 0 & 1 \end{pmatrix}$$

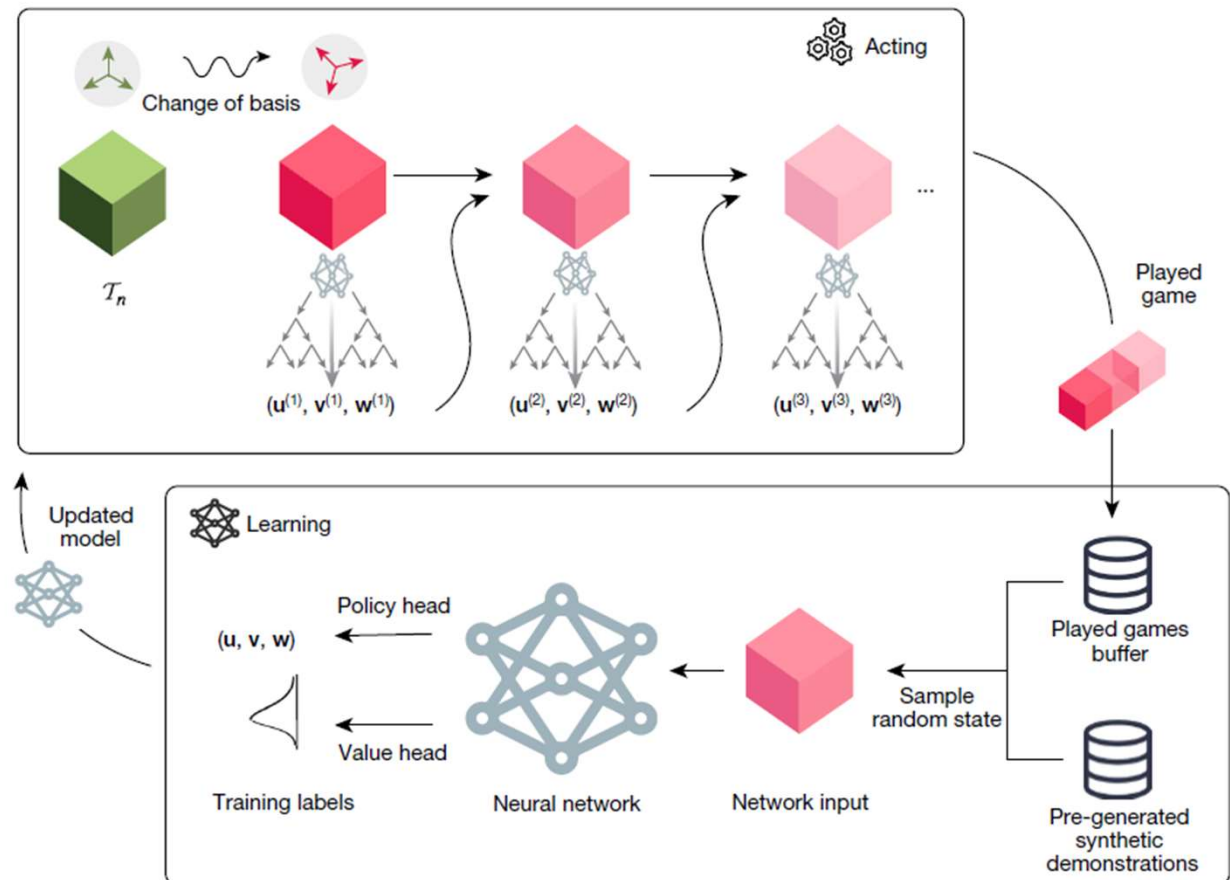
$$\mathbf{w} = \begin{pmatrix} 1 & 0 & 0 & 1 & -1 & 0 & 1 \\ 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & -1 & 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$

How to reward the agent

- Encourage finding shortest path: Each step is given -1 reward
- Encourage finding path closer to goal state: Terminal step given reward based on the complexity of the terminal tensor: $-\gamma(S_{R_{\text{limit}}})$
- u, v, w constrained to be of the elements $\{-2, -1, 0, 1, 2\}$

Acting and Learning Process

- Use Monte Carlo Tree Search to sample state space
- Policy Network tells us how often to choose an action
- Value Network gives the value of the current state


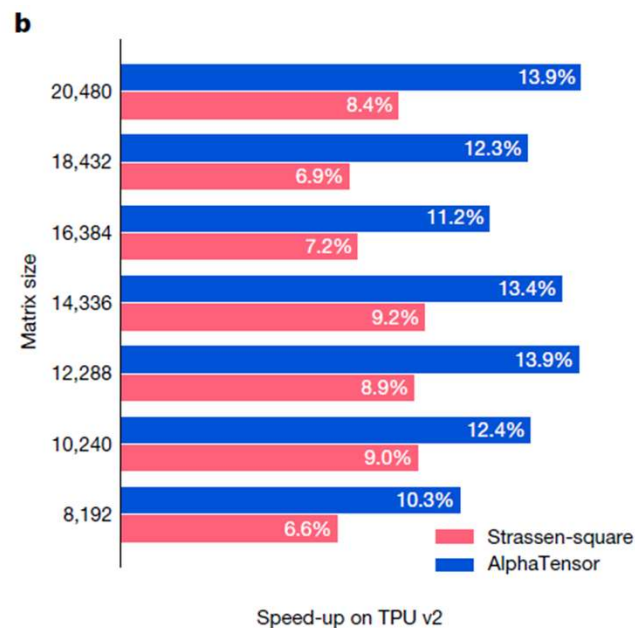
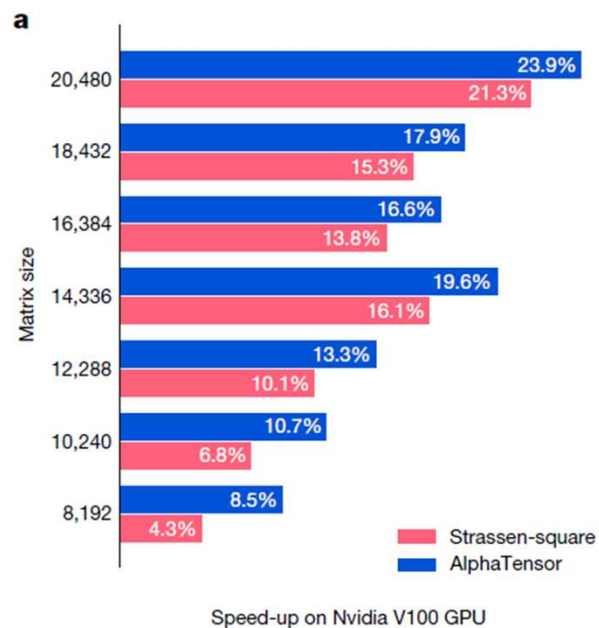


Additional Data Augmentation

- Synthetic demonstrations – Firstly, generate a custom goal tensor over R steps, then play the game using the known decomposition to give more “success” cases
 - Quite similar to Hindsight Experience Replay
- Change of Basis for Tensor Multiplication
 - Increase diversity of known situations
- Swapping a move with the last move (as game is order invariant)
 - Learning more from the same experience

Optimising for Hardware


- Adding a runtime-based reward at terminal state
 - Longer the runtime, the more negative this reward



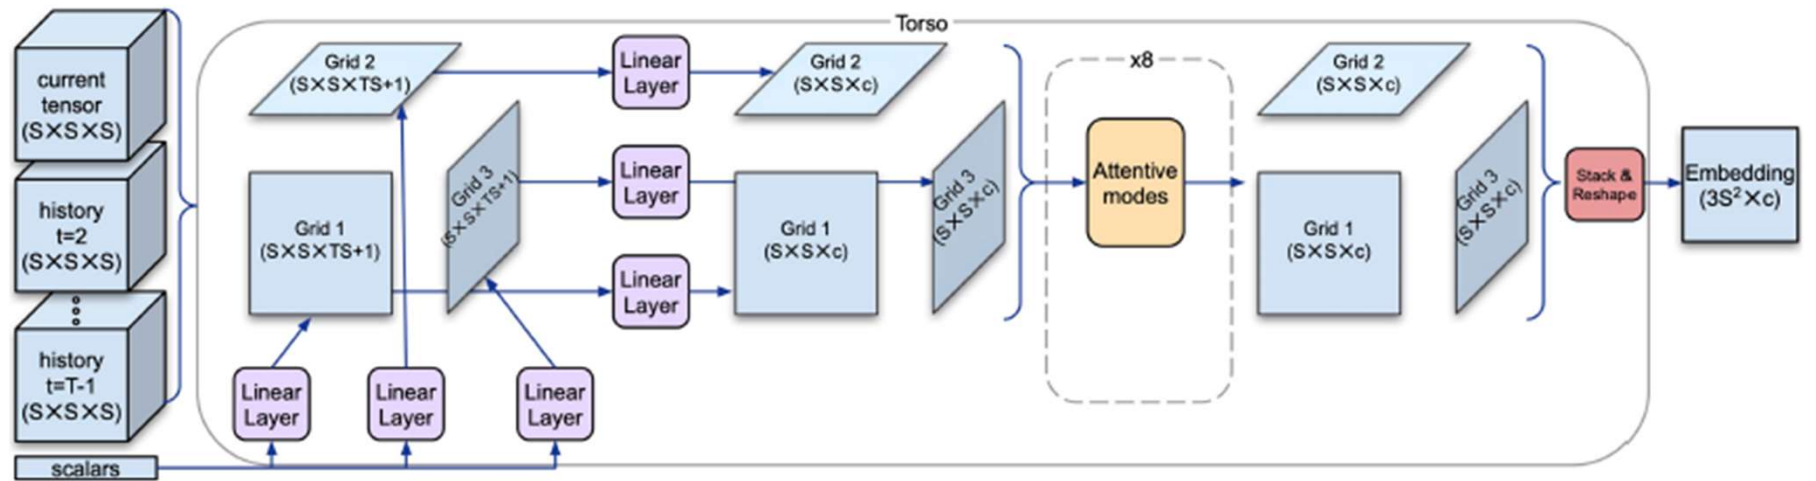
Training Time

- Learning:
 - TPU v3
 - 64 TPU cores
 - Train for 600,000 iterations
- Acting:
 - TPU v4
 - 1600 actors
- Training converges in 1 week

Network Architecture

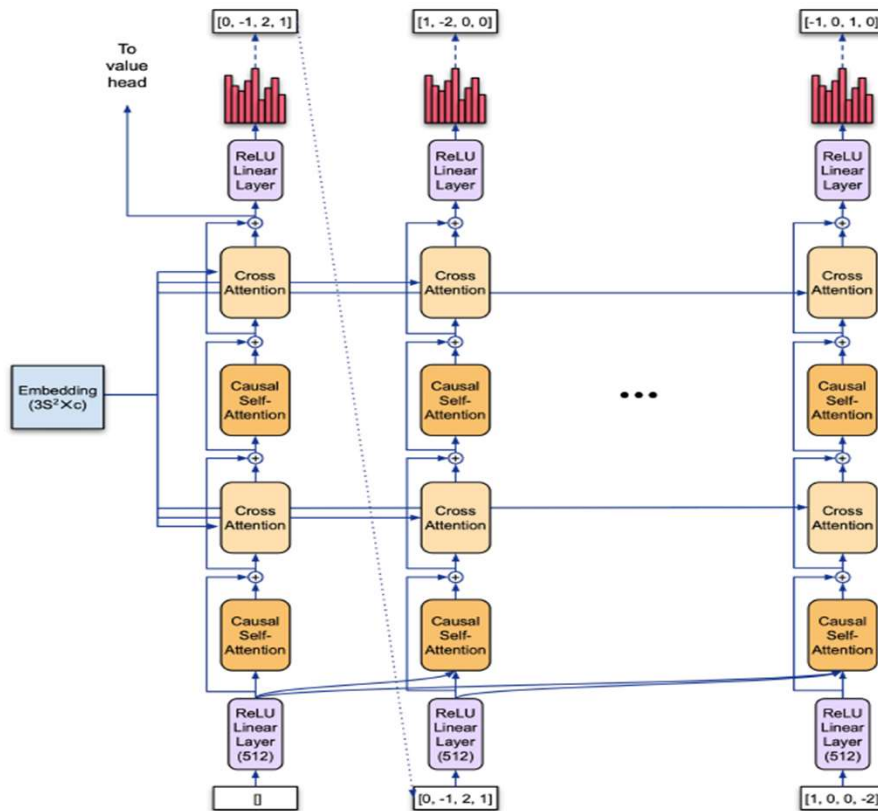


Torso




Policy and Value Heads

- Extensive use of Cross Attention



(b) Policy head diagram. Blocks in the same row share the learning weights.



(c) Value head diagram.

Thoughts

- A lot of inductive biases for matrix multiplication built into the system
 - Could help speed up learning
- Reward not just only win/loss/draw like AlphaZero, but more informative which is based on closeness to the goal state (and also compute time)
 - Similar to KataGo incorporating more rewards such as ownership loss and score belief loss to help the network focus on erroneous predictions
 - Easier to credit assign than just the final game outcome
- Attention mechanisms can help with credit assignment

Questions to Ponder

- Can we add more reward (e.g. intrinsic reward) to make the learning process more efficient?
- Can we sample the stored buffer better (e.g. hippocampal replay) to make the learning process more efficient?
- Could there be some form of algorithmic discovery whereby a sequence of actions lead to the elimination of one multiplication step, and we can apply this recursively?
- Is there a way to simplify TensorGame even more, by perhaps just considering fewer possibilities?
- How do we apply this to arbitrary one-player games?

FBHHRBNRSSHK-Algorithm

Manuel Kauers and Jakob Moosbauer

Overview

- Human-based multiplication of matrices
- Achieves 95 multiplications (instead of 96 found by AlphaTensor) for multiplication of 5×5 matrices
- Used the scheme of AlphaTensor
- Apply recursively:
 - Have a schema of multiplications
 - Applying a sequence of transformations to eliminate one multiplication

Problem Set-up

3. MULTIPLICATION OF 5×5 MATRICES

Let

$$A = \begin{pmatrix} a_{1,1} & a_{1,2} & a_{1,3} & a_{1,4} & a_{1,5} \\ a_{2,1} & a_{2,2} & a_{2,3} & a_{2,4} & a_{2,5} \\ a_{3,1} & a_{3,2} & a_{3,3} & a_{3,4} & a_{3,5} \\ a_{4,1} & a_{4,2} & a_{4,3} & a_{4,4} & a_{4,5} \\ a_{5,1} & a_{5,2} & a_{5,3} & a_{5,4} & a_{5,5} \end{pmatrix} \text{ and } B = \begin{pmatrix} b_{1,1} & b_{1,2} & b_{1,3} & b_{1,4} & b_{1,5} \\ b_{2,1} & b_{2,2} & b_{2,3} & b_{2,4} & b_{2,5} \\ b_{3,1} & b_{3,2} & b_{3,3} & b_{3,4} & b_{3,5} \\ b_{4,1} & b_{4,2} & b_{4,3} & b_{4,4} & b_{4,5} \\ b_{5,1} & b_{5,2} & b_{5,3} & b_{5,4} & b_{5,5} \end{pmatrix},$$

and

$$C = \begin{pmatrix} c_{1,1} & c_{1,2} & c_{1,3} & c_{1,4} & c_{1,5} \\ c_{2,1} & c_{2,2} & c_{2,3} & c_{2,4} & c_{2,5} \\ c_{3,1} & c_{3,2} & c_{3,3} & c_{3,4} & c_{3,5} \\ c_{4,1} & c_{4,2} & c_{4,3} & c_{4,4} & c_{4,5} \\ c_{5,1} & c_{5,2} & c_{5,3} & c_{5,4} & c_{5,5} \end{pmatrix} := AB$$

be matrices with coefficients in a ring R of characteristic 2.

Inclusion-Exclusion Matrices

We can then compute the entries of C from the entries of A and B by first setting

$$m_k = \left(\sum_{i=1}^5 \sum_{j=1}^5 \alpha_{i,j}^{(k)} a_{i,j} \right) \left(\sum_{i=1}^5 \sum_{j=1}^5 \beta_{i,j}^{(k)} b_{i,j} \right)$$

for $k = 1, \dots, 95$ and then

$$c_{i,j} = \sum_{k=1}^{95} \gamma_{i,j}^{(k)} m_k$$

for $i, j = 1, \dots, 5$, where the $\alpha_{i,j}^{(k)}, \beta_{i,j}^{(k)}, \gamma_{i,j}^{(k)}$ are as follows:

k	$((\alpha_{i,j}^{(k)}))_{i,j=1}^5$	$((\beta_{i,j}^{(k)}))_{i,j=1}^5$	$((\gamma_{i,j}^{(k)}))_{i,j=1}^5$
1	$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{pmatrix}$	$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{pmatrix}$	$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{pmatrix}$